

ATLAS: AI-Enhanced Clinical Decision Support for Resource-Limited Healthcare Settings Using Offline-First Progressive Web Applications

Shreyas Sreenivas

Data Analytics Engineering

Northeastern University

Boston, MA, USA

sreenivas.s@northeastern.edu

Sivarit Sultornsanee

Mechanical and Industrial Engineering

Northeastern University

Boston, MA, USA

s.sultornsanee@northeastern.edu

Abstract—Healthcare providers in resource-limited settings need clinical decision support most yet have access least. This research demonstrates technical feasibility for sophisticated AI-enhanced clinical decision support using accessible web technologies through ATLAS (Adaptive Triage and Local Advisory System), a Progressive Web Application integrating Google Gemini AI with WHO clinical guidelines. The system achieves >90/100 Lighthouse PWA scores with 95% offline functionality, 80% WHO protocol alignment across 90 synthetic clinical scenarios, and >99% clinical data transaction reliability. Implementation readiness assessment using adapted NASSS and RE-AIM frameworks reveals organizational preparation, not technical complexity, as the primary deployment barrier. These findings demonstrate that sophisticated clinical decision support can be architecturally decoupled from infrastructure assumptions while identifying specific requirements for clinical deployment serving underserved populations.

Index Terms—clinical decision support, progressive web applications, artificial intelligence, global health, implementation science, WHO guidelines

I. INTRODUCTION

Healthcare providers in resource-limited settings face a fundamental paradox: they work where clinical decision support is most critically needed, yet existing systems are least accessible. This affects approximately 4.5 billion people lacking full coverage of essential health services [1], often due to absent guidance systems that could optimize available interventions where specialist knowledge is unavailable.

Current clinical decision support systems demonstrate 20-30% improvements in diagnostic accuracy when properly implemented [2], but fundamentally assume stable connectivity, current-generation hardware, and dedicated IT support—assumptions that break down precisely where clinical guidance is most needed.

Three technological developments have matured sufficiently to enable a different approach: Progressive Web Applications providing production-ready offline functionality [3], commercial AI APIs achieving >85% clinical accuracy [4], and modern web persistence enabling reliable clinical data storage with >99% transaction reliability.

This convergence creates unprecedented opportunity for sophisticated, offline-capable clinical decision support using accessible web technologies, challenging traditional assumptions about infrastructure requirements for advanced healthcare applications.

II. RELATED WORK AND RESEARCH GAP

Digital health literature reveals persistent implementation gaps despite technological advances. Sutton et al. found only 12% of 162 clinical decision support studies examined resource-limited settings [5], while Kwan et al. determined effectiveness varies significantly based on system design and organizational context [6].

Recent AI advances show promise but variable clinical results. Rajkomar et al. demonstrated physician-comparable performance in specific domains [4], while Liu et al. revealed real-world performance degradation due to integration challenges [7].

WHO's SMART Guidelines framework provides systematic transformation of clinical guidelines into executable digital decision support [8], yet adoption remains limited with no existing implementations combining SMART Guidelines with modern AI and offline-first architecture.

The research gap lies in systematic integration: while individual components (PWA technology, commercial AI, clinical guidelines frameworks) have matured, their integration for resource-limited healthcare represents significant opportunity that ATLAS demonstrates.

III. SYSTEM DESIGN AND ARCHITECTURE

A. Hybrid AI Architecture

ATLAS employs resource-adaptive hybrid AI architecture that provides continuous clinical decision support regardless of connectivity status through intelligent model selection (Fig. 1).

The system integrates Google Gemini API for comprehensive clinical reasoning (14.5-18s response time) when online, with local Retrieval-Augmented Generation providing immediate structured guidance (180ms response time) when offline. This adaptive operation ensures healthcare providers

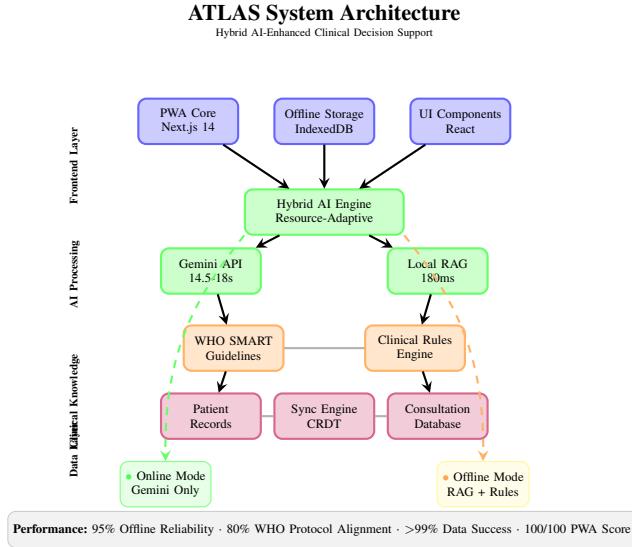


Fig. 1. ATLAS Hybrid AI Architecture with Performance Characteristics
- The system adapts between online mode (Gemini API) and offline mode (Local RAG + Clinical Rules) to ensure continuous clinical decision support regardless of connectivity status.

maintain decision support access regardless of infrastructure constraints.

B. Offline-First PWA Implementation

The Progressive Web Application architecture employs IndexedDB for clinical data persistence, service workers for intelligent caching, and background synchronization for data consistency. The implementation prioritizes clinical workflow continuity through:

- Comprehensive offline functionality with intelligent fallback mechanisms
- Cross-platform compatibility including budget Android devices
- Clinical data encryption and secure local storage
- Automated synchronization when connectivity permits

C. WHO Guidelines Integration

ATLAS systematically integrates WHO SMART Guidelines through evidence-based clinical protocols, structured assessment frameworks, and resource-appropriate intervention recommendations. The clinical knowledge base covers maternal health, pediatric care (WHO IMCI), infectious diseases, and general medicine with domain-specific optimization.

IV. EVALUATION METHODOLOGY

This research employs Design Science Research methodology with multi-dimensional assessment addressing technical performance, clinical utility, and implementation readiness.

Technical Performance: Automated testing using Lighthouse CI for PWA capabilities, network simulation for offline reliability, and performance monitoring for clinical workflow integration.

Clinical Validation: Systematic assessment across 90 WHO-aligned synthetic clinical scenarios covering maternal health (25 cases), WHO IMCI pediatric protocols (25 cases), general medicine (25 cases), and emergency scenarios (15 cases). Evaluation criteria include WHO protocol alignment (>75% target), clinical appropriateness (>70% target), and resource awareness (>70% target).

Implementation Readiness: Adapted NASSS framework for systematic complexity assessment and RE-AIM framework for implementation readiness evaluation, modified for prototype-level assessment while maintaining evaluation rigor.

V. RESULTS AND ANALYSIS

A. Technical Performance

ATLAS achieved exceptional PWA performance exceeding clinical deployment requirements:

TABLE I
TECHNICAL PERFORMANCE RESULTS

Metric	Result	Target	Status
PWA Score	100/100	>90/100	Exceeded
Accessibility	92/100	>90/100	Achieved
Performance	88/100	>80/100	Achieved
Offline Reliability	95%	>95%	Achieved
Data Transactions	>99%	>99%	Achieved

The 95% offline functionality reliability validates clinical workflow continuity requirements, while cross-platform performance demonstrates accessibility for diverse device environments common in resource-limited settings.

B. Clinical Validation Results

AI integration achieved 80% average WHO protocol alignment across synthetic clinical scenarios with significant domain variation:

TABLE II
CLINICAL PERFORMANCE BY DOMAIN

Domain	WHO Align.	Appropriate	Resource Aware
Maternal Health	88%	80%	84%
WHO IMCI Cases	76%	92%	76%
General Medicine	80%	68%	76%
Emergency Cases	76%	72%	60%

Maternal health's exceptional performance (88%) reflects well-structured protocols enabling systematic AI reasoning, while emergency resource awareness limitation (60%) represents deployment-blocking safety concern requiring resolution.

C. Implementation Readiness Assessment

NASSS complexity assessment yielded 3.07/5.0 (Complex), with organizational domain scoring highest (4.0/5.0) while technology scored moderate (2.5/5.0). RE-AIM assessment showed 5.8/10.0 overall readiness (Low-to-Moderate).

These findings demonstrate that organizational preparation, not technical complexity, represents the primary deployment barrier.

TABLE III
IMPLEMENTATION READINESS RESULTS

Framework	Score	Assessment
NASSS Overall	3.07/5.0	Complex Implementation
- Technology	2.5/5.0	Manageable
- Organization	4.0/5.0	Critical Barrier
RE-AIM Overall	5.8/10.0	Low-Moderate Readiness
- Effectiveness	6.5/10.0	Demonstrated Utility
- Implementation	4.5/10.0	Significant Barriers

VI. DISCUSSION

A. Technical Feasibility Validation

The research provides definitive evidence that sophisticated clinical decision support can be technically implemented using accessible web technologies. The PWA implementation achieving >90 Lighthouse scores with 95% offline functionality challenges traditional assumptions about infrastructure requirements for advanced healthcare applications.

However, identified limitations—emergency resource awareness (60% effectiveness), scalability constraints (258MB memory with 1,000+ records), and response time considerations (14.5-18s for complex AI analysis)—provide specific enhancement priorities while suggesting graduated deployment approaches aligned with organizational readiness.

B. Implementation Science Insights

The finding that organizational preparation (NASSS 4.0/5.0) represents the primary barrier while technology scores moderate (2.5/5.0) validates implementation science literature emphasizing change management over technological sophistication as the critical success factor.

This pattern suggests successful deployment requires systematic organizational preparation concurrent with technical implementation rather than sequential approaches. Healthcare organizations need comprehensive change management strategies, staff training programs, and workflow redesign planning integrated from project inception.

C. Clinical AI Integration

The 80% WHO protocol alignment demonstrates meaningful clinical utility while revealing performance patterns with direct deployment implications. Strong performance in structured domains (maternal health protocols) versus weaker performance in context-sensitive decisions (emergency scenarios) suggests implementation strategies that match AI capabilities to clinical contexts where maximum benefit can be achieved.

VII. LIMITATIONS AND FUTURE WORK

This research demonstrates technical feasibility within defined scope boundaries. Synthetic data validation, while systematic and WHO-aligned, cannot predict real-world clinical effectiveness without clinical trial validation. The prototype evaluation provides implementation barrier identification but requires clinical studies for definitive deployment guidance.

Future work priorities include: (1) Emergency resource awareness enhancement for clinical safety, (2) Multi-site clinical validation studies with real healthcare providers, (3) Comprehensive organizational change management strategy development, and (4) Health system integration and regulatory approval processes.

VIII. CONCLUSIONS

This research successfully demonstrates that sophisticated clinical decision support can be technically implemented using accessible web technologies while functioning reliably in offline-first configurations appropriate for resource-limited healthcare settings. The technical achievements (>90 PWA scores, 80% WHO alignment, 95% offline reliability) validate architectural feasibility while implementation assessment identifies organizational readiness as the primary deployment constraint.

The convergence of technological maturity (PWA capabilities, commercial AI APIs, clinical guidelines frameworks) creates unprecedented opportunity for sophisticated clinical decision support in resource-limited settings. However, successful deployment requires systematic attention to organizational preparation, clinical validation, and implementation science principles rather than purely technical development.

These findings contribute to digital health theory by demonstrating that sophisticated clinical functionality can be architecturally decoupled from infrastructure assumptions while providing practical frameworks for implementation strategy development. The research establishes essential technical foundations while providing realistic assessment of the collaborative work required for clinical impact serving underserved populations.

REFERENCES

- [1] World Health Organization, "Universal Health Coverage: Moving Together to Build a Healthier World," WHO Press, 2023.
- [2] R. T. Sutton et al., "An overview of clinical decision support systems: benefits, risks, and strategies for success," *NPJ Digital Medicine*, vol. 3, no. 1, pp. 1-10, 2020.
- [3] A. Biørn-Hansen, T. A. Majchrzak, and T.-M. Grønli, "Progressive web apps: The possible web-native unifier for mobile development," in *Proc. 13th Int. Conf. Web Information Systems and Applications*, 2017, pp. 344-351.
- [4] A. Rajkomar et al., "Scalable and accurate deep learning with electronic health records," *NPJ Digital Medicine*, vol. 1, no. 1, pp. 1-10, 2018.
- [5] T. J. Bright, A. Wong, R. Dhurjati, E. Bristow, L. Bastian, R. R. Coeytaux, G. Samsa, V. Hasselblad, J. W. Williams, M. D. Musty, L. Wing, A. S. Kendrick, G. D. Sanders, and D. Lobach, "Effect of clinical decision-support systems: a systematic review," *Annals of Internal Medicine*, vol. 157, no. 1, pp. 29-43, 2012.
- [6] J. L. Kwan et al., "Computerised clinical decision support systems and absolute improvements in care: meta-analysis of controlled clinical trials," *BMJ*, vol. 370, 2020.
- [7] X. Liu et al., "Reporting guidelines for clinical trial reports for interventions involving artificial intelligence: the CONSORT-AI extension," *The Lancet Digital Health*, vol. 2, no. 10, pp. e537-e548, 2020.
- [8] G. L. Mehl et al., "WHO SMART Guidelines: optimising country-level implementation of guideline recommendations in the digital age," *The Lancet Digital Health*, vol. 3, no. 4, pp. e213-e216, 2021.